

UNIVERSITÀ DEGLI STUDI DI MILANO
FACULTY OF SCIENCE AND TECHNOLOGY

DEPARTMENT OF COMPUTER SCIENCE



Master in
Computer Science

STATISTICAL METHODS FOR MACHINE LEARNING
FINAL REPORT ON NEURAL NETWORKS FOR THE
BINARY CLASSIFICATION

Teacher: Prof. Nicolò Cesa-Bianchi

Final report written by:
Samuele Simone
Matr. Nr. 11910A

ACADEMIC YEAR 2022-2023

I/We declare that this material, which I/We now submit for assessment, is entirely my/our own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my/our work. I/We understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me/us or any other person for assessment on this or any other course of study.

Contents

Index	ii
1 Introduction	1
2 Dataset	2
2.1 Preprocessing	4
3 Neural Network Models	6
3.1 K-fold cross validation (K = 5)	6
3.2 Multi-Layer Perceptron (MLP)	7
3.2.1 First MLP Model	7
3.3 Convolutional Neural Network (CNN)	8
3.4 Deep Residual learning (ResNet50)	8
4 Experiments	9
References	11

Chapter 1

Introduction

The purpose of this project was to discover and work on different neural network models in order to solve an image classification problem. Specifically the recognition of chihuahuas from muffins, which because of the similarities, turns out to be a non-trivial task for a machine to perform. In the following chapters we will go into detail about the various models and through experiments and graphs figure out which model will work best to solve this task.

Chapter 2

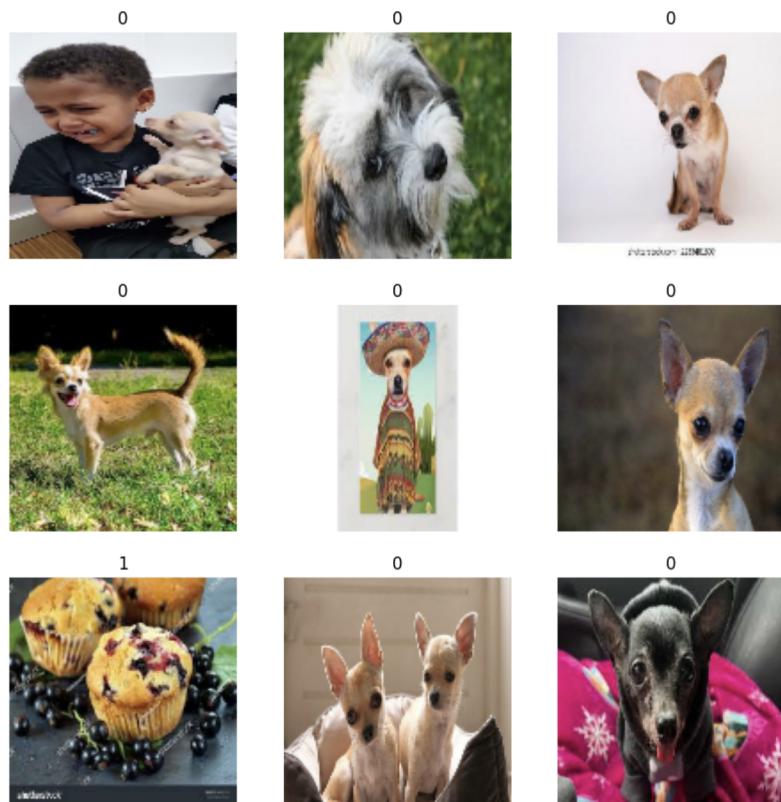
Dataset

The dataset under consideration is called "Muffin vs Chihuahua", taken from Kaggle [1]. It consists of about 6000 images taken from Google Images. Duplicate images have been removed. The first operation that was performed was to download the dataset through the Kaggle API and divide the dataset into training set and test set. After that, thanks to the Tensorflow library and specifically Keras, the images were loaded so that the neural networks could be trained. Here the code for the train_images loading using `tf.keras.utils.image_dataset_from_directory`

```
1 train_images = tf.keras.utils.image_dataset_from_directory(  
2     train_folder,  
3     labels="inferred",  
4     label_mode="int",  
5     class_names=None,  
6     color_mode="rgb",  
7     batch_size=32,  
8     image_size=(128, 128),  
9     shuffle=True,  
10    seed=None,  
11    validation_split=None,  
12    subset=None,  
13    interpolation="bilinear",  
14    follow_links=False,  
15    crop_to_aspect_ratio=False,  
16 )
```

As you can see there are different properties such as `label_mode = int`, so labels as integers. Or `color_mode = "rgb"` in which the images are working in this red blue green space. By making these parameters explicit, it is possible to obtain a more custom configuration that suits our needs. The same process was applied for the test_images. To give an idea to the reader, here a sample of images found within the dataset: As you can see muffin label is 0 and chihuahua label is 1.

Figure 1: A sample of images found within the dataset



2.1 Preprocessing

Certainly when dealing with data it is good to always make some improvement such as cleaning the data so that you have better results. In fact, the dataset was delivered without having duplicate images, which already represents a first step of preprocessing. In addition, I applied the **Data augmentation** [2] that is a technique in machine learning used to reduce overfitting when training a machine learning model, by training models on several slightly-modified copies of existing data. With the help of Keras, it is possible to achieve date augmentation in the following way:

```

1 data_augmentation = keras.Sequential(
2     [
3         layers.RandomFlip("horizontal"),
4         layers.RandomRotation(0.1),
5     ]
6 )

```

Here is graphically what the transformation looks like: In order to complete the

Figure 2: Data augmentation applied in a specific image



preprocessing I applied the data augmentation transformation to the input and I

performed normalization of pixel values in the range [0, 255] to values in the range [0, 1], indeed this helps stabilize the training of the model. To conclude, this line `train_images = train_images.prefetch(tf.data.AUTOTUNE)` optimizes the data loading process during model training, allowing the model to work more efficiently and reducing the waiting time between data batches.

Chapter 3

Neural Network Models

In this chapter we are going to examine the models that have been made in order to solve the binary classification task. All models were addressed using the principle of K-Fold cross validation. In this case the $K = 5$.

3.1 K-fold cross validation ($K = 5$)

Cross-validation is a statistical method used to estimate the skill of machine learning models. The procedure has a single parameter called k that refers to the number of groups that a given data sample is to be split into. As such, the procedure is often called k -fold cross-validation. When a specific value for k is chosen, it may be used in place of k in the reference to the model such as $k=5$ becoming 5-fold cross-validation, the one used in this project. [3]. The general rules are the following:

1. Shuffle the dataset randomly
2. Split the dataset into k groups
3. For each group:
 - Take the group as a hold out or test data set
 - Take the remaining groups as a training data set
 - Fit a model on the training set and evaluate it on the test set
 - Retain the evaluation score and discard the model
4. Summarize the skill of the model using the sample of model evaluation scores

Let's take a look at its mathematical formalization. Let S be our entire dataset [4]. We partition S in K subsets (also known as folds) S_1, \dots, S_K of size m/K each (assume for simplicity that K divides m). The K -fold CV estimate of $\mathbb{E}[\ell_{\mathcal{D}}(A)]$ on S , denoted by $\ell_S^{\text{cv}}(A)$, is then computed as follows: we run A on each training part S_{-i} of the folds $i = 1, \dots, K$ and obtain the predictors $h_1 = A(S_{-1}), \dots, h_K = A(S_{-K})$. We then compute the (rescaled) errors on the testing part of each fold,

$$\ell_{S_i}(h_i) = \frac{K}{m} \sum_{(\mathbf{x}, y) \in S_i} \ell(y, h_i(\mathbf{x}))$$

Finally, we compute the CV estimate by averaging these errors

$$\ell_S^{\text{cv}}(A) = \frac{1}{K} \sum_{i=1}^K \ell_{S_i}(h_i)$$

3.2 Multi-Layer Perceptron (MLP)

A multilayer perceptron (MLP) is a feedforward artificial neural network, consisting of fully connected neurons with a nonlinear kind of activation function, organized in at least three layers, notable for being able to distinguish data that is not linearly separable. [5] During the project, I developed 3 different MLP models so that I could conduct different experiments and understand the performance of them.

3.2.1 First MLP Model

Here are the MLP architecture:

- `inputs = tf.keras.Input(shape=input_shape)`: model's input with the `input_shape` size. In this case (128, 128, 3)
- `x = layers.Flatten()(inputs)`: Converts 2D input (an image) to a 1D vector
- `x = layers.Dense(256, activation='relu')(x)`: Fully connected layer (dense) with 256 neurons and ReLU activation
- `x = layers.Dropout(0.5)(x)`: This dropout layer introduces regularization into the network. The parameter 0.5 indicates that 50% of the neurons exiting this layer are randomly switched off during training.
- `x = layers.Dense(128, activation='relu')(x)`: Other fully connected layer with 128 neurons and ReLU activation.

- `x = layers.Dropout(0.5)(x)`: Another dropout layer for further regularization.
- `activation = 'sigmoid'`: Here the activation function for the last layer is specified. In your case, you are using 'sigmoid,' which is commonly used in binary classification problems where the output must be a probability between 0 and 1 for each class.
- `outputs = layers.Dense(num_classes, activation=activation)(x)`. `num_classes = 1`. This is the last layer of the model, which returns the output. Then these are the technical details of the setup:
 - `batch_size = 32`. Indicates how many data examples are processed together before updating the weights.
 - `optimizer='adam'`. The optimizer is the algorithm that adjusts the model weights during training to minimize the cost function. It is a stochastic gradient descent method.
 - `loss = 'binary_crossentropy'`. It specifies the cost function (or loss function) that will be used to evaluate how well the model is fitting the data during training. This cross-entropy loss is user for binary (0 or 1) classification applications.
 - `metrics=[zero_one_loss_func]`: this param specifies the evaluation metrics that will be calculated during model training. In this case I used a custom function called `zero_one_loss_func`. It requires two parameters: `y_true`, the real associated label and the `y_pred`, the ones predicted by the model. Then:
 - * Transforms predicted probabilities into binary labels (0 or 1)
 - * Compare the predicted binary labels with the actual labels
 - * Calculate the zero-one loss as the average of the errors.

3.3 Convolutional Neural Network (CNN)

3.4 Deep Residual learning (ResNet50)

Chapter 4

Experiments

Here the tables results for each experiment that I conducted:

Table 1: First MLP training results for each fold and epoch

Fold	Epoch	Loss	zero_one_loss_func	Val_Loss	Val_zero_one_loss_func
1	1	841.5640	0.4819	0.7778	0.4574
1	2	0.8091	0.4618	0.7777	0.4584
1	3	0.7503	0.4589	0.7795	0.4584
1	4	0.6898	0.4604	0.7803	0.4584
1	5	7.5421	0.4615	0.7684	0.4564
1	6	0.6897	0.4607	0.7778	0.4564
1	7	0.6899	0.4607	0.7784	0.4564
1	8	0.7288	0.4603	0.7786	0.4564
1	9	0.6897	0.4609	0.7786	0.4564
1	10	0.7136	0.4598	0.7782	0.4564
2	1	920.0229	0.4724	0.7597	0.4522
2	2	2.3125	0.4600	0.7584	0.4532
2	3	0.7153	0.4623	0.6895	0.4522
2	4	0.8337	0.4622	0.6886	0.4522
2	5	0.6901	0.4614	0.6891	0.4522
2	6	0.6899	0.4597	0.6890	0.4522
2	7	0.6898	0.4609	0.6890	0.4522
2	8	0.7029	0.4594	0.6890	0.4522
2	9	0.6902	0.4620	0.6890	0.4522
2	10	0.6903	0.4620	0.6890	0.4522
3	1	593.7564	0.4803	0.6951	0.4432

Continued on next page

Table 1 – Continued from previous page

Fold	Epoch	Loss	zero_one_loss_func	Val_Loss	Val_zero_one_loss_func
	2	3.0590	0.4624	0.6904	0.4432
	3	0.6901	0.4641	0.6895	0.4432
	4	0.6900	0.4641	0.6889	0.4432
	5	0.7200	0.4620	0.6885	0.4432
	6	0.6898	0.4618	0.6882	0.4432
	7	0.6895	0.4647	0.6881	0.4432
	8	0.6907	0.4638	0.6880	0.4432
	9	0.6897	0.4641	0.6880	0.4432
	10	0.6901	0.4635	0.6880	0.4432
	4	907.8840	0.4852	0.7944	0.4707
	2	2.0251	0.4560	0.7090	0.4718
	3	0.6919	0.4543	0.7087	0.4707
	4	0.6900	0.4563	0.7087	0.4707
	5	0.6895	0.4548	0.7088	0.4707
	6	0.6896	0.4570	0.7089	0.4707
	7	0.6918	0.4550	0.7001	0.4697
	8	0.6896	0.4575	0.7002	0.4697
	9	0.6894	0.4560	0.7003	0.4697
	10	0.6893	0.4550	0.7003	0.4697
	5	830.7402	0.5075	0.6967	0.4719
	2	1.6454	0.4563	0.6913	0.4708
	3	0.8036	0.4571	0.6909	0.4708
	4	0.6914	0.4576	0.6907	0.4708
	5	0.6896	0.4563	0.6907	0.4708
	6	0.6895	0.4553	0.6908	0.4708
	7	0.6895	0.4560	0.6908	0.4708
	8	0.6892	0.4571	0.6909	0.4708
	9	0.6893	0.4566	0.6910	0.4708
	10	0.6893	0.4563	0.6910	0.4708

Bibliography

- [1] Samuel Cortinhas. muffin-vs-chihuahua-image-classification. <https://www.kaggle.com/datasets/samuelcortinhas/muffin-vs-chihuahua-image-classification>, 2023.
- [2] Wikipedia. Data augmentation. https://en.wikipedia.org/wiki/Data_augmentation, 2023.
- [3] Jason Brownlee. A gentle introduction to k-fold cross-validation. <https://machinelearningmastery.com/k-fold-cross-validation/>, 2020.
- [4] Nicolò Cesa-Bianchi. Hyperparameter tuning and risk estimates. <https://cesa-bianchi.di.unimi.it/MSA/Notes/crossVal.pdf>, 2023.
- [5] Wikipedia. Multilayer perceptron. https://en.wikipedia.org/wiki/Multilayer_perceptron, 2023.